

Pipelining

Array multiplier implementation

$$-a_{N-1}2^{N-1} + \sum_{i=0}^{N-2} a_i 2^i$$
$$\begin{aligned} P &= \left(-b_{M-1}2^{M-1} + \sum_{j=0}^{M-2} b_j2^j \right) \cdot \left(-a_{N-1}2^{N-1} + \sum_{i=0}^{N-2} a_i2^i \right) \\ &= \sum_{i=0}^{N-2} \sum_{j=0}^{M-2} a_i b_j 2^{i+j} + a_{N-1} b_{M-1} 2^{M+N-2} - \left(\sum_{i=0}^{N-2} a_i b_{M-1} 2^{i+M-1} + \sum_{j=0}^{M-2} a_{N-1} b_j 2^{j+N-1} \right) \end{aligned}$$
$$P = \sum_{i=0}^4 \sum_{j=0}^4 a_i b_j 2^{i+j} + a_5 b_5 2^{10} - \left(\sum_{i=0}^4 a_i b_5 2^{i+5} + \sum_{j=0}^4 b_j a_5 2^{j+5} \right)$$

								b_5	b_4	b_3	b_2	b_1	b_0
								a_5	a_4	a_3	a_2	a_1	a_0
$\sum_{i=0}^4 \sum_{j=0}^4 a_i b_j 2^{i+j}$	[$b_4 a_0$	$b_3 a_0$	$b_2 a_0$	$b_1 a_0$	$b_0 a_0$
								$b_4 a_1$	$b_3 a_1$	$b_2 a_1$	$b_1 a_1$	$b_0 a_1$	
							$b_4 a_2$	$b_3 a_2$	$b_2 a_2$	$b_1 a_2$	$b_0 a_2$		
						$b_4 a_3$	$b_3 a_3$	$b_2 a_3$	$b_1 a_3$	$b_0 a_3$			
						$b_4 a_4$	$b_3 a_4$	$b_2 a_4$	$b_1 a_4$	$b_0 a_4$			
$a_5 b_5 2^{10}$			$b_5 a_5$										
$-\sum_{i=0}^4 a_i b_5 2^{i+5}$		1	1	$\overline{a_4 b_5}$	$\overline{a_3 b_5}$	$\overline{a_2 b_5}$	$\overline{a_1 b_5}$	$\overline{a_0 b_5}$	1	1	1	1	1
													1
$-\sum_{j=0}^4 b_j a_5 2^{j+5}$		1	1	$\overline{b_4 a_5}$	$\overline{b_3 a_5}$	$\overline{b_2 a_5}$	$\overline{b_1 a_5}$	$\overline{b_0 a_5}$	1	1	1	1	1
													1
		P_{11}	P_{10}	P_9	P_8	P_7	P_6	P_5	P_4	P_3	P_2	P_1	P_0

This is the standard Baugh-Wooley 2's complement multiplication algorithm which results in the array multiplier in Figure 3. Note that there is no need for sign extension of the partial products. The Baugh-Wooley array multiplier usually is implemented using carry-save adder modules as this reduces the delay through the array. However, in FPGA, ripple carry adders are fast because of the fast carry chain logic within a logic slice, so it is easier to use standard adders (as shown in Figure 3). The parameterised Verilog code for implementing N-bit x N-bit array multiplier is given in NTUlearn.

							b_5	b_4	b_3	b_2	b_1	b_0
							a_5	a_4	a_3	a_2	a_1	a_0
						1	b_5a_0	b_4a_0	b_3a_0	b_2a_0	b_1a_0	b_0a_0
					b_5a_1		b_4a_1	b_3a_1	b_2a_1	b_1a_1	b_0a_1	
				b_5a_2	b_4a_2		b_3a_2	b_2a_2	b_1a_2	b_0a_2		
			b_5a_3	b_4a_3	b_3a_3		b_2a_3	b_1a_3	b_0a_3			
		b_5a_4	b_4a_4	b_3a_4	b_2a_4		b_1a_4	b_0a_4				
1	b_5a_5	b_4a_5	b_3a_5	b_2a_5	b_1a_5	b_0a_5						
P_{11}	P_{10}	P_9	P_8	P_7	P_6	P_5	P_4	P_3	P_2	P_1	P_0	

Figure 2. Two's complement multiplication (rearranged) with 6 partial products

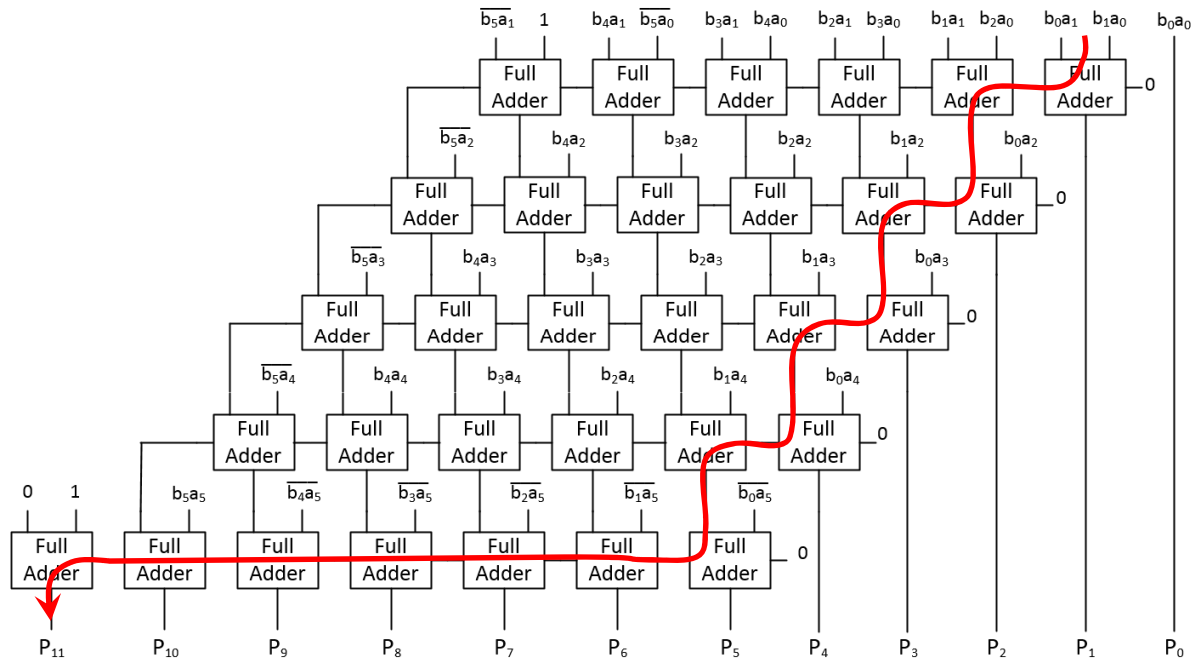


Figure 3. The 6-bit x 6-bit 2's complement array multiplier implementation

Task 1

Create a new project called Lab3 in Xilinx Vivado, targeting the Basys 3 board (Artix-7 xc7a35tcbg236-1 FPGA). **Remember to use a local drive location NOT a network drive.** Download the lab3.zip file from NTUlearn and unzip it into the Lab3 directory.

Open the file `array_mult_parameterised.v` in a text editor. This code will generate any arbitrary N-bit x N-bit multiplier. It is currently set to an 8-bit x 8-bit multiplier. But remember Verilog allows you to override the parameters in an upper level module. To generate a 6-bit multiplier, we could use:

```
array_mult #(.SIZE(6)) uut (.a(a), .b(b), .P(P));
```

Spot the generation of the partial products and then the intermediate sums and the final result. However, this code is a bit difficult to understand, and so we will use a non-parameterised 6-bit version for this lab. Next, open `array_mult_6Bit.v` in a text editor. This code directly matches to Figure 3. Compare the two implementations. Note, both use the same resources and produce identical hardware. Close the files.

Then add the Verilog design files (`AM_top.v` and `array_mult_6Bit.v`) to the design (select **+** OR **Add Sources** from **PROJECT MANAGER** in the **Flow Navigator** window). Also add the constraints file (`Lab3.xdc`). The `Lab3.xdc` file contains the pin constraints for the design. Change the clock constraint to a 9ns clock (as shown below) and **synthesise the design**.

```
create_clock -add -name sys_clk_pin -period 9 -waveform {0 4.5} [get_ports clk]
```

Select **Reports** in the bottom window, then open the `synth_1_synth_synthesis_report_0` report and verify that only the resource expected has been used (2 x 6-bit registers and 1 x 12-bit register and 5 x 7-bit adders).

Then **Open Synthesized Design** from **SYNTHESIS** in the **Flow Navigator** window. Then select **Report Timing Summary**. Click **OK** at the pop-up window. Then in the bottom window, select “Timing” and verify that all of the **user specified timing constraints are met**. Note that the *Worst Negative Slack* (WNS) and *Worst hold Slack* (WHS) should be fractions of a ns, meaning that we could not operate the circuit (the clock) much faster. That is, not much below 8.5ns (9ns is the current value). **IMPORTANT: Use a sniping tool to keep a copy of the “Design Timing Summary” window for assessment purposes.**

Task 2

Next we will simulate the design. **Close the SYNTHESIZED DESIGN window** and then select **Add Sources**. Select **Add or create simulation sources**, then **Next**. Select **Add files**, and then select *AM_top_tb.v* and click **OK**. Click **Finish**. The test bench instantiates the top module and provides a 10ns clock and selected multiplier inputs.

Run the Behavioural Simulation. Select *uut* from **Scope** in the **SIMULATION** window. In the **Objects** window select *a_r* and *b_r* and move them to the **Name** section of the waveform window (the black part). This will add the registered multiplier inputs (e.g. aligned to the clock) to the simulation. Move *result[11:0]* to the bottom (below the other signals). Then **Restart** and **Run All** to redo the simulation with the new signals. Change the radix to signed decimal and check the multiplier gives the correct results. The waveform window should look like Figure 4. Note the registered signals are aligned to the clock and the result is delayed by 1 clock cycle.

Task 3

Now we should pipeline the multiplier design. First, examine the block diagram of the multiplier shown in Figure 3. Note that the critical path is from the inputs to the full adder (FA) block at the top right (b_0a_1 and b_1a_0) to the output from the FA block at the bottom left (P_{11}), and is 15 FA delays. So we need to partition the adder somewhere such that the critical path in the two partitions stays roughly the same. From Figure 3, inserting a pipeline register between the 3rd and 4th adder blocks, gives a critical path of 10 FA delays for the top partition (9 for the bottom). The pipelined design is as shown in Figure 5 (and should use the same signals).

Close the SIMULATION window. Remove *array_mult_6Bit.v* from the project (right click on it and then select “Remove File from project”). Then add *multA.v* and *multB.v* to the design. Open *multA.v* and check that it only uses the first three partial products and that the 10-bit *P* is assigned as the sum of *PP3* and *Isum2* (plus some lower terms). Next open *multB.v*. It uses the original *a* and *b* inputs as well as the 10-bit product (*Pin*) generated in *multA.v* and produces a 12-bit result (*P*). Verify that the two new modules match those shown in Figure 5.

Then modify *AM_top* to remove the *array_mult_6Bit* instantiation and instantiate the two new multiplier modules (remember that *multB.v* has an extra input (*Pin*)). You will also need to add appropriate pipeline registers of the correct bit widths. The pipeline register outputs (inputs to the two modules) should be declared as type *reg* while the module outputs will be declared as wires. Next modify the always @ (posedge *clk*) block.

Open the constraints file (*Lab.xdc*) and edit the clock constraint, and change to:

```
create_clock -add -name sys_clk_pin -period 5.5 -waveform {0 2.75} [get_ports clk]
```

Now, synthesize the pipelined design. There are several additional warnings in the design relating to unconnected ports. These can be ignored as a quick check of the *multA.v* and *multB.v* Verilog code (in the “generate partial products” section) will show that *multA* does not use *a[5:4]* and *multB* does not use *a[3:0]*. Again “Open (the) Synthesised Design” and select “Report Timing Summary”. Then **OK**.

Then select “Timing” and verify that the **user specified timing constraints** are met. Note in the unpipelined design we used a 9ns clock (111.1MHz), whereas here we are now using a 5.5ns clock (181.8MHz). This represents a 64% increase in clock frequency due to pipelining. A significant increase. **IMPORTANT: Again, use a sniping tool to keep a copy of the “Design Timing Summary” window for assessment purposes.**

Simulate the design. You should be able to use the same test bench from Task 2. At the simulation window, add the 5 pipeline register outputs (*a_r1*, *a_r2*, *b_r1*, *b_r2* and *Pin*) to the simulation window. Move *result[11:0]* to the bottom. Then add a divider and add *Pa* and *Pb* at the bottom. **Restart** and **Run All** to redo the simulation with the new signals. Verify the results are correct. Why is *result* just after reset non-zero?

Inform your lab supervisor once you reach this point

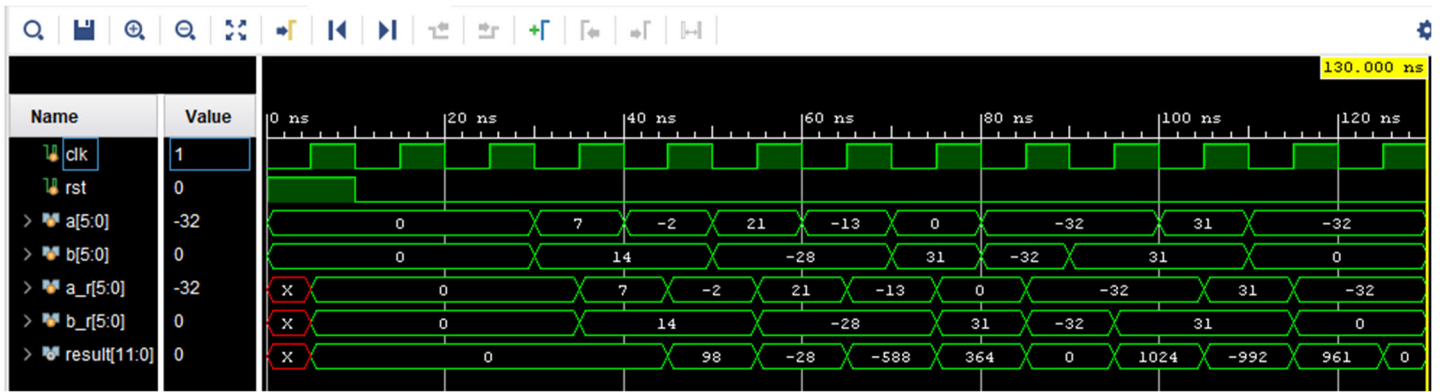


Figure 4. Simulation waveform for non-pipelined multiplier

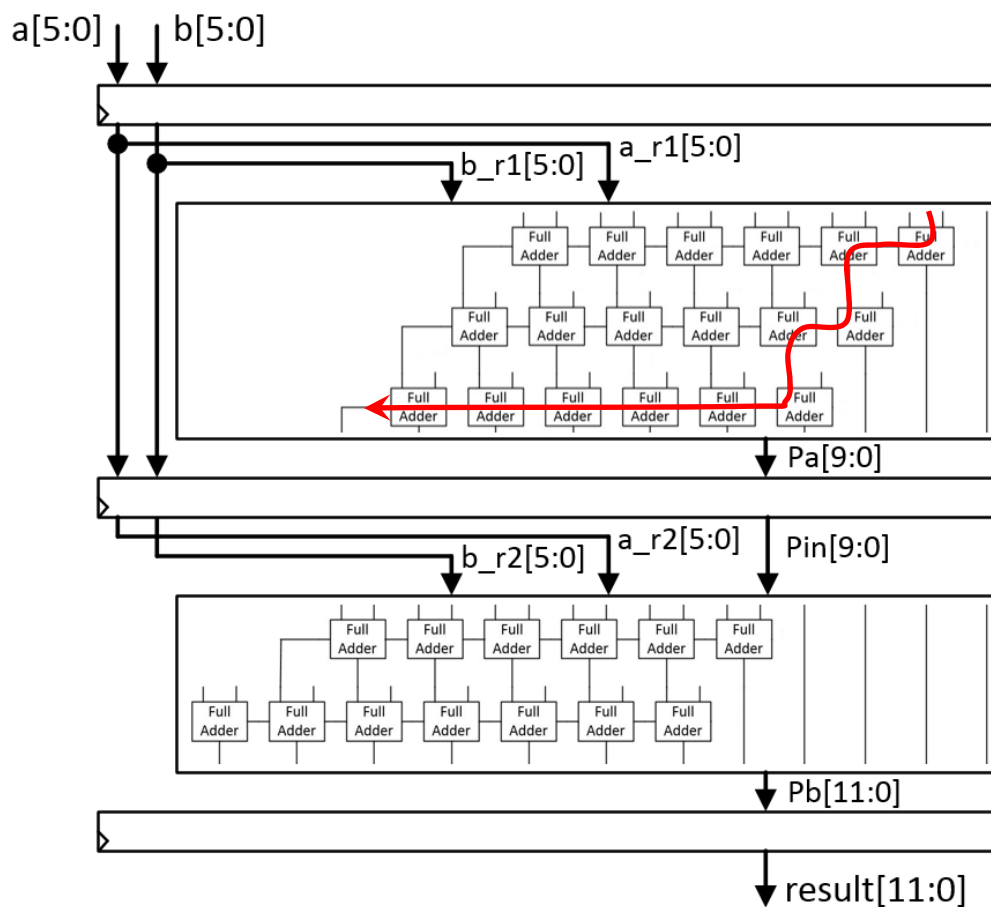


Figure 5. The block diagram of the pipelined multiplier

Note: The following warnings can be safely ignored

Synthesis (Task 1):

WARNING: [Constraints 18-5210] No constraints selected for write.

Synthesis (Task 3):

WARNING: [Synth 8-3331] design multB has unconnected port a[3] (5 more like this)

WARNING: [Constraints 18-5210] No constraints selected for write.